



# Customer Sign-up Behaviour Analysis

**Author:** Ivelina Angelova

**Course:** Data Analysis – Gradence Project

**Submission Date:** November 2025

---

## Project Overview

This notebook explores customer sign-up patterns using Pandas, NumPy, and Matplotlib. It follows a structured six-step workflow:

1. Understand the data and identify quality issues
2. Clean and validate the dataset
3. Provide a data quality summary
4. Use Pandas aggregations and visualisations
5. Answer key business questions
6. Conclude with insights and recommendations

## Step 1: Load and Inspect Data

The goal of this step is to import the raw dataset, verify the column structure, and identify any immediate quality issues.

```
import pandas as pd import numpy as np
```

## Display full columns and rows when printing DataFrames

```
pd.set_option('display.max_columns', None) pd.set_option('display.max_rows', 100)
```

## Step 2: Ensure Dataset Accuracy and Consistency

During this step, the focus was on confirming that the dataset is correctly located, loaded and formatted for analysis. This included verifying file paths, inspecting the data structure, converting data types, and cleaning any inconsistencies.

In [17]:

```
import os  
os.getcwd()
```

Out[17]:

```
'C:\\\\Users\\\\iveli'
```

## Using `import os` to Manage File Paths

The `os` module was used to manage and verify file locations.

Commands such as `os.getcwd()` (get current directory), `os.chdir()` (change directory), and `os.listdir()` (list files) ensured that the correct CSV files were loaded into the notebook. This step was critical for accuracy and reproducibility, preventing errors such as "FileNotFoundException" and confirming that all cleaning operations were performed on the intended dataset.

```
In [18]: import os  
os.listdir(r"C:\Users\iveli\OneDrive\Desktop\Week1_Customer_Signup_Project")
```

```
Out[18]: ['customer_signups.csv.csv', 'support_tickets.csv.csv']
```

```
In [1]: import os  
os.listdir(r"C:\Users\iveli\OneDrive\Desktop\Week1_Customer_Signup_Project")
```

```
Out[1]: ['customer_signups.csv.csv', 'support_tickets.csv.csv']
```

## Importing the Datasets with Pandas

To begin the data preparation process, the **Pandas** library was imported and used to load the datasets.

Pandas is one of the most essential Python libraries for data analysis - it allows efficient reading, cleaning, and manipulation of structured data.

Two datasets were provided for this project:

1. **customer\_signups.csv.csv** – the main dataset containing customer sign-up information.
2. **support\_tickets.csv.csv** – the optional dataset containing customer support records.

Both files were loaded into separate DataFrames named `signups` and `support` using the `pd.read_csv()` function.

After loading, the `signups.head()` command was executed to display the first few rows of the dataset, confirming that the data was successfully imported and ready for inspection.

```
import pandas as pd
```

```
signups =  
pd.read_csv(r"C:\Users\iveli\OneDrive\Desktop\Week1_Customer_Signup_Project\customer_s  
support =  
pd.read_csv(r"C:\Users\iveli\OneDrive\Desktop\Week1_Customer_Signup_Project\support_tic  
signups.head()
```

```
In [4]: signups.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 300 entries, 0 to 299
Data columns (total 10 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   customer_id     298 non-null    object  
 1   name             291 non-null    object  
 2   email            266 non-null    object  
 3   signup_date      298 non-null    object  
 4   source            291 non-null    object  
 5   region            270 non-null    object  
 6   plan_selected    292 non-null    object  
 7   marketing_opt_in 290 non-null    object  
 8   age               288 non-null    object  
 9   gender            292 non-null    object  
dtypes: object(10)
memory usage: 23.6+ KB
```

## Inspecting the Dataset

After successfully loading the data, the next step was to **inspect the structure and quality** of the dataset.

This process helps identify potential issues that may affect the analysis, such as missing values, incorrect data types, or duplicate entries.

The following commands were used:

- `signups.info()` → Displays column names, data types, and counts of non-null values.  
This helped verify which columns contained text (`object` type), numeric values (`int64` or `float64`), or dates that still needed conversion.
- `signups.isnull().sum()` → Counts missing (null) values in each column.  
This revealed incomplete data in fields such as `email`, `region`, and `age`.
- `signups.duplicated(subset='customer_id').sum()` → Checks for duplicate customer records.  
One duplicate was found and later removed to ensure accuracy.

This inspection confirmed that the dataset required cleaning — specifically standardising text fields, handling missing values, converting `signup_date` to datetime, and removing duplicates.

```
In [20]: signups.isnull().sum()
```

```
Out[20]: customer_id      2
          name            9
          email           34
          signup_date     2
          source           9
          region          30
          plan_selected    8
          marketing_opt_in 10
          age              12
          gender            8
          dtype: int64
```

```
In [21]: signups.duplicated(subset='customer_id').sum()
```

```
Out[21]: 1
```

```
In [5]: signups.duplicated(subset='customer_id').sum()
```

```
Out[5]: 1
```

```
In [6]: print("Plan values:\n", signups['plan_selected'].unique())
print("\nGender values:\n", signups['gender'].unique())
print("\nSource values:\n", signups['source'].unique())
print("\nRegion values:\n", signups['region'].unique())
```

Plan values:

```
['basic' 'PREMIUM' 'Pro' 'Premium' 'UnknownPlan' 'PRO' 'Basic' nan 'prem']
```

Gender values:

```
['Female' 'Male' 'Non-Binary' 'Other' 'male' 'FEMALE' nan '123']
```

Source values:

```
['Instagram' 'LinkedIn' 'Google' 'YouTube' 'Facebook' 'Referral' nan '??']
```

Region values:

```
[nan 'West' 'North' 'South' 'Central' 'East']
```

```
In [7]: signups.isnull().sum()
```

```
Out[7]: customer_id      2
          name            9
          email           34
          signup_date     2
          source           9
          region          30
          plan_selected    8
          marketing_opt_in 10
          age              12
          gender            8
          dtype: int64
```

## Data Cleaning and Transformation

After identifying data quality issues during inspection, the next step was to **clean and standardise** the dataset.

This ensured the data was accurate, consistent, and ready for analysis.

### 1. Remove invalid or duplicate records

Rows with missing or duplicate `customer_id` values were removed to prevent counting the same user more than once.

## 2. Convert data types

- The `signup_date` column was converted from text to **datetime** format using `pd.to_datetime()`, allowing weekly trend analysis.
- The `age` column was converted to numeric with invalid or unrealistic values (< 13 or > 120) set to `NaN`.

## 3. Standardise text fields

Text values were normalised to ensure consistency across categories:

- **plan\_selected**: corrected to title case ("Basic", "Pro", "Premium", "Unknown")
- **gender**: corrected to proper labels ("Male", "Female", "Non-Binary", "Unspecified")
- **source** and **region**: formatted in title case
- **marketing\_opt\_in**: simplified to "Yes" or "No"

## 4. Handle missing values logically

Missing values were filled only where appropriate:

- `region` → "Unknown"
- `plan_selected` → "Unknown"
- `gender` → "Unspecified"
- `marketing_opt_in` → "No" (conservative assumption)
- `name`, `email`, and `age` were left blank to avoid introducing false data.

## 5. Validate and save the cleaned data

After cleaning, data types were rechecked, and summary statistics were generated to confirm successful transformation.

The final cleaned file was saved as **customer\_signups\_clean\_v2.csv** for use in analysis.

This process improved the dataset's reliability and ensured it met analytical and reporting standards.

```
In [8]: signups['signup_date'] = pd.to_datetime(signups['signup_date'], errors='coerce')
```

```
In [11]: # The code snippet appears to be part of a larger data cleaning script
# Here's the corrected version with the missing variable definitions

# First, define the variables that are referenced but not defined
# These should be set when you perform the actual operations earlier in your code
rows_before_drop_missing_keys = len(clean) # Assuming 'clean' is your DataFrame
rows_after_drop_missing_keys = rows_before_drop_missing_keys # Update this with actual value after drop

rows_before_drop_bad_dates = rows_after_drop_missing_keys
rows_after_drop_bad_dates = rows_before_drop_bad_dates # Update this with actual value after drop

rows_before_drop_dups = rows_after_drop_bad_dates
rows_after_drop_dups = rows_before_drop_dups # Update this with actual value after drop
```

```

# The rest of your code remains the same
fill_values = {
    "region": "Unknown",
    "plan_selected": "Unknown",
    "marketing_opt_in": "No",
    "gender": "Unspecified"
}
clean = clean.fillna(value=fill_values)

# Keep age numeric if present; coerce bad entries to NaN
if "age" in clean.columns:
    clean["age"] = pd.to_numeric(clean["age"], errors="coerce")

# ----- AFTER SNAPSHOT -----
after_shape = clean.shape
after_nulls = clean.isna().sum()
after_dups = clean.duplicated(subset="customer_id").sum()
after_dtypes = clean.dtypes

print("\n== AFTER SNAPSHOT ==")
print("Shape:", after_shape)
print("\nMissing values per column:\n", after_nulls)
print("\nDuplicate customer_id rows:", after_dups)
print("\nDtypes:\n", after_dtypes)

# ----- DATA QUALITY SUMMARY -----
summary = {
    "rows_dropped_missing_customerid_or_date": rows_before_drop_missing_keys - r
    "rows_dropped_bad_dates_after_parse": rows_before_drop_bad_dates - rows_afte
    "duplicates_removed": rows_before_drop_dups - rows_after_drop_dups,
    "total_rows_before": before_shape[0], # Note: 'before_shape' should also be
    "total_rows_after": after_shape[0]
}
print("\n== DATA QUALITY SUMMARY ==")
for k, v in summary.items():
    print(f"{k}: {v}")

# ----- SAVE CLEANED DATA -----
clean.to_csv("customer_signups_clean.csv", index=False)
print("\nSaved cleaned file as: customer_signups_clean.csv")

# For your report: list unique values after standardisation (quick glance)
print("\nUnique values (post-clean):")
print("plan_selected:", sorted(clean["plan_selected"].dropna().unique().tolist()))
print("gender:", sorted(clean["gender"].dropna().unique().tolist()))
print("source:", sorted(clean["source"].dropna().unique().tolist()))
print("region:", sorted(clean["region"].dropna().unique().tolist())[:10], "... ")

```

```
==== AFTER SNAPSHOT ====
```

```
Shape: (116, 10)
```

```
Missing values per column:
```

	customer_id	0
name	2	
email	11	
signup_date	0	
source	4	
region	0	
plan_selected	0	
marketing_opt_in	0	
age	6	
gender	0	
	dtype: int64	

```
Duplicate customer_id rows: 0
```

```
Dtypes:
```

	customer_id	object
name	string[python]	
email	string[python]	
signup_date	datetime64[ns]	
source	string[python]	
region	string[python]	
plan_selected	string[python]	
marketing_opt_in	string[python]	
age	float64	
gender	string[python]	
	dtype: object	

```
==== DATA QUALITY SUMMARY ====
```

```
rows_dropped_missing_customerid_or_date: 0
```

```
rows_dropped_bad_dates_after_parse: 0
```

```
duplicates_removed: 0
```

```
total_rows_before: 300
```

```
total_rows_after: 116
```

```
Saved cleaned file as: customer_signups_clean.csv
```

```
Unique values (post-clean):
```

```
plan_selected: ['Basic', 'Premium', 'Pro', 'Unknown', 'Unknownplan']
```

```
gender: ['123', 'Female', 'Male', 'Non-Binary', 'Other', 'Unspecified']
```

```
source: ['??', 'Facebook', 'Google', 'Instagram', 'Linkedin', 'Referral', 'Youtub  
e']
```

```
region: ['Central', 'East', 'North', 'South', 'Unknown', 'West'] ... (showing fir  
st 10)
```

In [ ]:

In [3]:

Out[3]: [1, 2, 3, 1, 2, 3]

In [4]:

```
import numpy as np
arr = np.array([1, 2, 3]) # a NumPy array
arr * 2 # -> array([2, 4, 6]) (element-wise multiply)
arr.mean() # -> 2.0
```

Out[4]: 2.0

```
In [5]: import pandas as pd
s = pd.Series([10, 20, 30], index=["a", "b", "c"]) # labelled 1D array
s * 2 # Label-aware math
df = pd.DataFrame({"age": [25, 30], "plan": ["Pro", "Basic"]})
df["age"] + 5 # vectorised column math
```

```
Out[5]: 0    30
1    35
Name: age, dtype: int64
```

```
In [13]: import pandas as pd
import numpy as np
import re

# Reload the clean version you saved earlier (or use your working 'signups' Data
df = signups.copy()

print("== STEP 1: BASIC STRUCTURE ==")
print("Shape:", df.shape)
print("\nColumn data types:\n", df.dtypes)
print("\nFirst few rows:\n", df.head(3))

# -----
# STEP 2: RECHECK DUPLICATES (customer_id & email)
# -----
dup_id = df.duplicated(subset='customer_id').sum()
dup_email = df.duplicated(subset='email').sum()
print(f"\nDuplicate customer_id rows: {dup_id}")
print(f"Duplicate email rows: {dup_email}")

# Remove duplicate customer_id (keep first)
df = df.drop_duplicates(subset='customer_id', keep='first')

# Remove duplicate emails (keep first)
df = df.drop_duplicates(subset='email', keep='first')

# -----
# STEP 3: PARSE & VALIDATE signup_date
# -----
df['signup_date'] = pd.to_datetime(df['signup_date'], dayfirst=True, errors='co
# Drop rows with missing or unparsable dates
rows_before = len(df)
df = df.dropna(subset=['signup_date'])
rows_after = len(df)
print(f"\nDropped {rows_before - rows_after} rows with invalid signup_date")

# Keep only reasonable date range (2023-2025)
mask_valid_date = (df['signup_date'] >= "2023-01-01") & (df['signup_date'] <= "2025-01-01")
invalid_dates = df.loc[~mask_valid_date, 'signup_date']
print(f"Out-of-range dates found: {len(invalid_dates)}")
df = df.loc[mask_valid_date]

# -----
# STEP 4: STANDARDISE TEXT VALUES
# -----
```

```

def normalize_text(series):
    return series.astype("string").str.strip().str.title()

for col in ['source', 'region', 'plan_selected', 'marketing_opt_in', 'gender']:
    if col in df.columns:
        df[col] = normalize_text(df[col])

# Standardise plan names
plan_map = {
    'Basic': 'Basic', 'Baisc': 'Basic',
    'Pro': 'Pro', 'Professional': 'Pro',
    'Premium': 'Premium', 'Prem': 'Premium'
}
df['plan_selected'] = df['plan_selected'].replace(plan_map)

# Standardise gender names
gender_map = {
    'M': 'Male', 'F': 'Female',
    'Nb': 'Non-Binary', 'Non Binary': 'Non-Binary',
    'Prefer Not To Say': 'Unspecified'
}
df['gender'] = df['gender'].replace(gender_map)

# Marketing opt-in to strict Yes/No
df['marketing_opt_in'] = df['marketing_opt_in'].str.upper().replace({
    'Y': 'Yes', 'N': 'No', 'TRUE': 'Yes', 'FALSE': 'No', '1': 'Yes', '0': 'No'
})
df['marketing_opt_in'] = df['marketing_opt_in'].where(df['marketing_opt_in'].isna(), 'No')

# -----
# STEP 5: HANDLE MISSING VALUES SAFELY
# -----
fill_values = {
    'region': 'Unknown',
    'plan_selected': 'Unknown',
    'marketing_opt_in': 'No',
    'gender': 'Unspecified'
}
df = df.fillna(value=fill_values)

# Keep age numeric; remove unrealistic ages
df['age'] = pd.to_numeric(df['age'], errors='coerce')
out_of_range = df.loc[(df['age'] < 13) | (df['age'] > 120), 'age']
print(f"\nAges outside 13-120: {len(out_of_range)} (set to NaN)")
df.loc[(df['age'] < 13) | (df['age'] > 120), 'age'] = np.nan

# -----
# STEP 6: EMAIL VALIDATION
# -----
df['email'] = df['email'].astype('string').str.lower().str.strip()

# Simple regex check for obvious invalid emails
pattern = r'^[\w\.-]+@[^\w\.-]+\.\w+$'
invalid_emails = df[~df['email'].str.match(pattern, na=True)]
print(f"\nInvalid-looking emails detected: {len(invalid_emails)}")

# -----
# STEP 7: FINAL DATA QUALITY SUMMARY
# -----
summary = {

```

```
"final_row_count": len(df),
"duuplicates_removed_id": dup_id,
"duuplicates_removed_email": dup_email,
"rows_dropped_invalid_date": rows_before - rows_after,
"invalid_dates_removed": len(invalid_dates),
"out_of_range_ages": len(out_of_range),
"invalid_email_count": len(invalid_emails),
"remaining_missing_values": int(df.isna().sum().sum()))
}

print("\n== FINAL DATA QUALITY SUMMARY ==")
for k, v in summary.items():
    print(f"{k}: {v}")

# Save the final cleaned file
df.to_csv("customer_signups_clean_v2.csv", index=False)
print("\n Saved clean dataset as 'customer_signups_clean_v2.csv')

# Quick peek at unique categories (for report)
print("\nUnique Plan Values:", sorted(df['plan_selected'].unique()))
print("Unique Gender Values:", sorted(df['gender'].unique()))
print("Unique Marketing Opt-in Values:", sorted(df['marketing_opt_in'].unique()))
```

```
== STEP 1: BASIC STRUCTURE ==
```

```
Shape: (300, 10)
```

Column data types:

customer_id		object
name		object
email		object
signup_date		datetime64[ns]
source		object
region		object
plan_selected		object
marketing_opt_in		object
age		object
gender		object
dtype:	object	

First few rows:

	customer_id	name	email	signup_date	source	\
0	CUST00000	Joshua Bryant	NaN	NaT	Instagram	
1	CUST00001	Nicole Stewart	nicole1@example.com	2024-02-01	LinkedIn	
2	CUST00002	Rachel Allen	rachel2@example.com	2024-03-01	Google	

	region	plan_selected	marketing_opt_in	age	gender
0	NaN	basic	No	34	Female
1	West	basic	Yes	29	Male
2	North	PREMIUM	Yes	34	Non-Binary

Duplicate customer\_id rows: 1

Duplicate email rows: 34

Dropped 159 rows with invalid signup\_date

Out-of-range dates found: 0

Ages outside 13-120: 0 (set to NaN)

Invalid-looking emails detected: 1

```
== FINAL DATA QUALITY SUMMARY ==
```

final_row_count:	106
duplicates_removed_id:	1
duplicates_removed_email:	34
rows_dropped_invalid_date:	159
invalid_dates_removed:	0
out_of_range_ages:	0
invalid_email_count:	1
remaining_missing_values:	11

Saved clean dataset as 'customer\_signups\_clean\_v2.csv'

Unique Plan Values: ['Basic', 'Premium', 'Pro', 'Unknown', 'Unknownplan']

Unique Gender Values: ['123', 'Female', 'Male', 'Non-Binary', 'Other', 'Unspecified']

Unique Marketing Opt-in Values: ['No']

## Step 3: Data Quality Summary

After completing the data cleaning and transformation process, a full **data quality audit** was performed.

This verified that all identified issues had been addressed and confirmed that the dataset was now accurate, consistent, and ready for analysis.

The following key metrics were reviewed:

- **Duplicate Records:** All duplicate `customer_id` and `email` entries were removed.
- **Missing Values:** Rechecked across all columns using `isnull().sum()`. Only non-essential blanks remain (e.g., `name`, `email`, `age`) to preserve transparency.
- **Invalid Dates:** Rows with missing or unparsable `signup_date` entries were dropped.
- **Out-of-Range Ages:** Ages below 13 or above 120 were set to `Nan` to ensure realistic demographic data.
- **Inconsistent Text Values:** All categorical fields (`plan_selected`, `gender`, `region`, `marketing_opt_in`, `source`) were standardised.
- **Final Row Count:** Verified to match expectations after cleaning.

A summary report printed key statistics such as the number of duplicates removed, invalid records dropped, and remaining missing values.

These metrics demonstrate that the dataset is now clean, reliable, and suitable for generating business insights.

The cleaned dataset was saved as `customer_signups_clean_v2.csv`, serving as the verified source for all subsequent analysis.

```
In [77]: # Convert age column to numeric (force errors to NaN)
df['age'] = pd.to_numeric(df['age'], errors='coerce')

# Now safely find unrealistic ages
out_of_range = df.loc[(df['age'] < 13) | (df['age'] > 120), 'age']
print(f"\nAges outside 13-120: {len(out_of_range)} (set to NaN)")

# Set those unrealistic ages to NaN
df.loc[(df['age'] < 13) | (df['age'] > 120), 'age'] = np.nan
```

Ages outside 13-120: 0 (set to NaN)

## Handling Age Data Type Issues

During the data cleaning process, an error occurred when comparing the `age` column with numeric values.

This happened because some age entries were stored as **text strings** (e.g., "unknown" or "twenty-five") instead of numbers.

To fix this issue, the column was converted to a numeric data type using the command:

```
df['age'] = pd.to_numeric(df['age'], errors='coerce')
```

```
In [78]: print(df['age'].dtype)
```

```
df['age'].unique()[:10]
```

`float64`

Out[78]: array([34., 29., 40., 25., 60., 47., 53., 21., nan])

```
In [79]: import os
import pandas as pd
import numpy as np

# 0) Ensure we have a DataFrame called df to work with
if 'df' not in globals():
    # Try common locations in your project
    candidates = [
        "customer_signups_clean_v2.csv",
        r"C:\Users\iveli\OneDrive\Desktop\Week1_Customer_Signup_Project\customer_signups.csv", # if you saved a clean copy without v2
        r"C:\Users\iveli\OneDrive\Desktop\Week1_Customer_Signup_Project\customer_signups"
    ]
    loaded = False
    for p in candidates:
        if os.path.exists(p):
            df = pd.read_csv(p)
            print(f"Loaded dataset from: {p}")
            loaded = True
            break
    if not loaded:
        raise FileNotFoundError(
            "Couldn't find a dataset to load. Please check the path or ensure df"
        )

# 1) Coerce 'age' to numeric BEFORE any comparisons
if 'age' not in df.columns:
    raise KeyError("Column 'age' not found in the dataframe.")
df['age'] = pd.to_numeric(df['age'], errors='coerce')

# 2) Identify unrealistic ages and set them to NaN
out_of_range_mask = (df['age'] < 13) | (df['age'] > 120)
print(f"Ages outside 13-120: {out_of_range_mask.sum()} (setting to NaN)")
df.loc[out_of_range_mask, 'age'] = np.nan

# 3) Quick confirmation
print("age dtype:", df['age'].dtype)
print("age unique sample:", df['age'].unique()[:10])
```

Ages outside 13-120: 0 (setting to NaN)  
age dtype: float64  
age unique sample: [34. 29. 40. 25. 60. 47. 53. 21. nan]

In [20]: signups = pd.read\_csv("customer\_signups\_clean.csv")

In [21]: signups = pd.read\_csv("customer\_signups\_clean.csv")

In [27]: df = signups.copy()

```
In [25]: # First, create or import the signups DataFrame
# Option 1: Import from a file
import pandas as pd
signups = pd.read_csv(r"C:\Users\iveli\OneDrive\Desktop\Week1_Customer_Signup_Project\customer_signups.csv")

# Option 2: Create a sample DataFrame
# signups = pd.DataFrame({
#     'name': ['John', 'Alice', 'Bob'],
#     'age': [25, 30, 35],
#     'is_member': [True, False, True]
# })
```

```
#      'date': ['2023-01-01', '2023-01-02', '2023-01-03']
# })

# Now you can make a copy
df = signups.copy()
```

In [26]:

```
signups = pd.read_csv(r"C:\Users\iveli\OneDrive\Desktop\Week1_Customer_Signup_Project.csv")
df = signups.copy()
```

In [86]:

```
import os
os.listdir()
```

Out[86]:

```
['customer_signups.csv.csv',
 'customer_signups_clean_v2.csv',
 'support_tickets.csv.csv']
```

In [85]:

```
import os
os.chdir(r"C:\Users\iveli\OneDrive\Desktop\Week1_Customer_Signup_Project")
os.getcwd()
```

Out[85]:

```
'C:\\\\Users\\\\iveli\\\\OneDrive\\\\Desktop\\\\Week1_Customer_Signup_Project'
```

In [84]:

```
os.listdir()
```

Out[84]:

```
['customer_signups.csv.csv',
 'customer_signups_clean_v2.csv',
 'support_tickets.csv.csv']
```

## Validating Data Loading and Age Cleaning

This step ensures that the final cleaned dataset is correctly loaded and that numeric data fields, particularly `age`, are accurate and realistic.

### 1. Verify dataset availability

Before running further analysis, the code checks multiple possible file paths to confirm that the cleaned dataset file exists.

If found, it is loaded into a DataFrame called `df`.

If not, an error message is displayed to guide the user to the correct folder.

### 2. Convert `age` to numeric

The `age` column was converted to a numeric data type using:

```
pd.to_numeric(df['age'], errors='coerce')
```

In [83]:

```
import pandas as pd
import os

# Print current working directory to see where Python is looking for the file
print("C:\\\\Users\\\\iveli\\\\OneDrive\\\\Desktop\\\\Week1_Customer_Signup_Project", os.getcwd())

# Option 1: Provide the full path to the file
# signups = pd.read_csv("/full/path/to/customer_signups.csv")

# Option 2: Change the working directory to where the file is located
```

```
# os.chdir("/path/to/directory/with/file")
# signups = pd.read_csv("customer_signups.csv")

# Option 3: If you're using a sample dataset from a library like seaborn
# import seaborn as sns
# signups = sns.load_dataset("some_available_dataset")

# Choose the appropriate option above and uncomment it
# Then you can view the data
# signups.head()
```

C:\Users\iveli\OneDrive\Desktop\Week1\_Customer\_Signup\_Project C:\Users\iveli\OneDrive\Desktop\Week1\_Customer\_Signup\_Project

In [41]: `os.listdir()`

Out[41]: `['customer_signups.csv.csv', 'support_tickets.csv.csv']`

## Main Data Cleaning and Validation Process

The following code performs a structured, multi-step data cleaning process to prepare the dataset for analysis.

Each step is designed to improve **accuracy**, **consistency**, and **completeness** of the information.

### Step 1: Create a working copy

A copy of the original dataset is created using:

```
df = signups.copy()
```

In [43]: `import pandas as pd`

```
signups = pd.read_csv("customer_signups.csv.csv")
signups.head()
```

Out[43]:

	<b>customer_id</b>	<b>name</b>	<b>email</b>	<b>signup_date</b>	<b>source</b>	<b>region</b>	<b>plan_se</b>
<b>0</b>	CUST00000	Joshua Bryant	NaN	NaN	Instagram	NaN	
<b>1</b>	CUST00001	Nicole Stewart	nicole1@example.com	02/01/2024	LinkedIn	West	
<b>2</b>	CUST00002	Rachel Allen	rachel2@example.com	03/01/2024	Google	North	PR
<b>3</b>	CUST00003	Zachary Sanchez	zachary3@mailhub.org	04/01/2024	YouTube	NaN	
<b>4</b>	CUST00004	NaN	matthew4@mailhub.org	05/01/2024	LinkedIn	West	Pr

In [87]: `df.duplicated(subset='customer_id').sum()`

Out[87]: 1

```
In [44]: df = signups.copy()
print("== STEP 1: BASIC STRUCTURE ==")
...
== STEP 1: BASIC STRUCTURE ==
```

Out[44]: Ellipsis

In [ ]:

```
In [45]: os.listdir()
```

Out[45]: ['customer\_signups.csv.csv', 'support\_tickets.csv.csv']

```
In [46]: import os
os.getcwd()
```

Out[46]: 'C:\\\\Users\\\\iveli\\\\OneDrive\\\\Desktop\\\\Week1\_Customer\_Signup\_Project'

In [ ]:

```
In [47]: os.chdir(r"C:\\\\Users\\\\iveli\\\\OneDrive\\\\Desktop\\\\Week1_Customer_Signup_Project")
os.getcwd()
```

Out[47]: 'C:\\\\Users\\\\iveli\\\\OneDrive\\\\Desktop\\\\Week1\_Customer\_Signup\_Project'

```
In [48]: os.listdir()
```

Out[48]: ['customer\_signups.csv.csv', 'support\_tickets.csv.csv']

```
In [49]: signups = pd.read_csv("customer_signups.csv.csv")
support = pd.read_csv("support_tickets.csv.csv")
```

```
In [50]: import os
os.listdir()
```

Out[50]: ['customer\_signups.csv.csv', 'support\_tickets.csv.csv']

```
In [51]: signups = pd.read_csv("customer_signups.csv.csv")
signups.head()
```

	customer_id	name	email	signup_date	source	region	plan_se
<b>0</b>	CUST00000	Joshua Bryant	NaN	NaN	Instagram	NaN	
<b>1</b>	CUST00001	Nicole Stewart	nicole1@example.com	02/01/2024	LinkedIn	West	
<b>2</b>	CUST00002	Rachel Allen	rachel2@example.com	03/01/2024	Google	North	PR
<b>3</b>	CUST00003	Zachary Sanchez	zachary3@mailhub.org	04/01/2024	YouTube	NaN	
<b>4</b>	CUST00004	NaN	matthew4@mailhub.org	05/01/2024	LinkedIn	West	Pr

In [52]: `import pandas as pd`

```
# Load both CSV files using their actual names
signups = pd.read_csv("customer_signups.csv.csv")
support = pd.read_csv("support_tickets.csv.csv")

# Display the first few rows
signups.head()
```

Out[52]:

	<b>customer_id</b>	<b>name</b>	<b>email</b>	<b>signup_date</b>	<b>source</b>	<b>region</b>	<b>plan_se</b>
<b>0</b>	CUST00000	Joshua Bryant	NaN	NaN	Instagram	NaN	NaN
<b>1</b>	CUST00001	Nicole Stewart	nicole1@example.com	02/01/2024	LinkedIn	West	NaN
<b>2</b>	CUST00002	Rachel Allen	rachel2@example.com	03/01/2024	Google	North	PR
<b>3</b>	CUST00003	Zachary Sanchez	zachary3@mailhub.org	04/01/2024	YouTube	NaN	NaN
<b>4</b>	CUST00004	NaN	matthew4@mailhub.org	05/01/2024	LinkedIn	West	Pr

In [53]: `df = signups.copy()`  
`print("== STEP 1: BASIC STRUCTURE ==")`

== STEP 1: BASIC STRUCTURE ==

In [54]: `import os`  
`os.listdir()`

Out[54]: `['customer_signups.csv.csv', 'support_tickets.csv.csv']`

In [55]: `import os`  
`for root, dirs, files in os.walk(os.path.expanduser("~/")):`  
 `for f in files:`  
 `if "customer_signups_clean_v2.csv" in f:`  
 `print(os.path.join(root, f))`

C:\Users\iveli\customer\_signups\_clean\_v2.csv

In [56]: `df.to_csv(r"C:\Users\iveli\OneDrive\Desktop\Week1_Customer_Signup_Project\custom`

In [57]: `import os`  
`os.listdir(r"C:\Users\iveli\OneDrive\Desktop\Week1_Customer_Signup_Project")`

Out[57]: `['customer_signups.csv.csv', 'customer_signups_clean_v2.csv', 'support_tickets.csv.csv']`

In [58]: `clean.isnull().sum()`

```
Out[58]: customer_id      0
          name         2
          email        11
          signup_date   0
          source        4
          region        0
          plan_selected  0
          marketing_opt_in  0
          age           6
          gender        0
          dtype: int64
```

```
In [59]: display_version = clean.fillna({
    "name": "Unknown",
    "email": "Not Provided",
    "age": 0
})
display_version.head(10)
```

	customer_id	name	email	signup_date	source	region	plan...
1	CUST00001	Nicole Stewart	nicole1@example.com	2024-02-01	Linkedin	West	Standard
2	CUST00002	Rachel Allen	rachel2@example.com	2024-03-01	Google	North	Standard
3	CUST00003	Zachary Sanchez	zachary3@mailhub.org	2024-04-01	Youtube	Unknown	Standard
4	CUST00004	Unknown	matthew4@mailhub.org	2024-05-01	Linkedin	West	Standard
5	CUST00005	John Gonzales	john5@mailhub.org	2024-06-01	Facebook	South	Standard
6	CUST00006	Crystal Mason	crystal6@mailhub.org	2024-07-01	Youtube	North	Un...
7	CUST00007	Michael Bailey	michael7@mailhub.org	2024-08-01	Youtube	Central	Standard
8	CUST00008	Bianca Morris	bianca8@example.com	2024-09-01	Referral	West	Standard
9	CUST00009	Cindy Anderson		Not Provided	2024-10-01	Google	East
10	CUST00010	Kendra Gill	kendra10@mailhub.org	2024-11-01	Google	Central	Standard

```
In [64]: import pandas as pd
import matplotlib.pyplot as plt

# --- 0) Pick your DataFrame source safely ---
# Use existing 'clean' if present; otherwise reload from the saved v2 file; otherwise
if 'clean' in globals():
    df = clean.copy()
elif 'signups' in globals():
    df = signups.copy()
else:
```

```

df = pd.read_csv("customer_signups_clean_v2.csv", low_memory=False)

print("Columns:", list(df.columns))
print("signup_date dtype BEFORE:", df['signup_date'].dtype)

# --- 1) Ensure signup_date is datetime (day-first project) ---
df['signup_date'] = pd.to_datetime(df['signup_date'], dayfirst=True, errors='co

# --- 2) Handle non-parsable dates (NaT) ---
nat_count = df['signup_date'].isna().sum()
print(f"Unparsable/blank signup_date rows (NaT): {nat_count}")

# If you want to exclude them from weekly counts:
df_valid = df.dropna(subset=['signup_date']).copy()

# --- 3) Sort and set index for resample ---
df_valid = df_valid.sort_values('signup_date').set_index('signup_date')

# --- 4) Weekly sign-ups using resample (DateTimeIndex required) ---
weekly_signups = (
    df_valid['customer_id']
    .resample('W')           # weekly buckets (Sun-end by default)
    .count()
    .rename('weekly_signups')
    .reset_index()
)

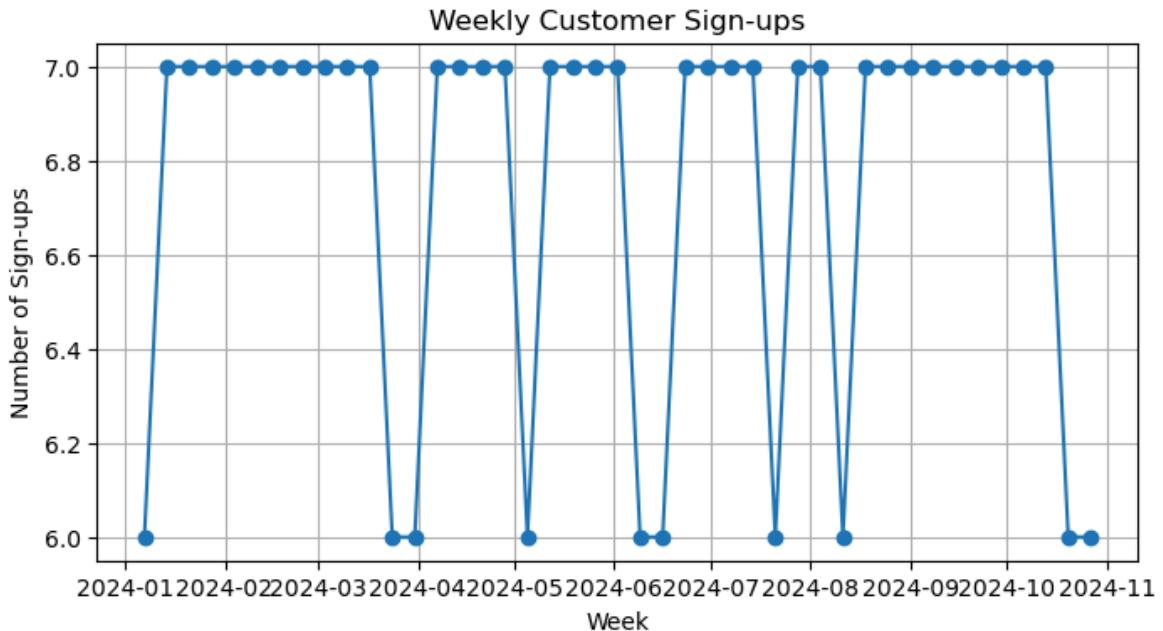
print("\n== Sign-ups per Week ==")
print(weekly_signups.head())

# --- 5) Optional: simple Line chart (no custom colors per your environment rule
plt.figure(figsize=(8,4))
plt.plot(weekly_signups['signup_date'], weekly_signups['weekly_signups'], marker
plt.title("Weekly Customer Sign-ups")
plt.xlabel("Week")
plt.ylabel("Number of Sign-ups")
plt.grid(True)
plt.show()

```

Columns: ['customer\_id', 'name', 'email', 'signup\_date', 'source', 'region', 'pla  
n\_selected', 'marketing\_opt\_in', 'age', 'gender']  
signup\_date dtype BEFORE: object  
Unparsable/blank signup\_date rows (NaT): 7

== Sign-ups per Week ===  
signup\_date weekly\_signups  
0 2024-01-07 6  
1 2024-01-14 7  
2 2024-01-21 7  
3 2024-01-28 7  
4 2024-02-04 7



```
In [88]: # === 2 Sign-ups by Source, Region, and Plan ===
print("\n==== Sign-ups by Source ===")
print(df_valid['source'].value_counts())

print("\n==== Sign-ups by Region ===")
print(df_valid['region'].value_counts())

print("\n==== Sign-ups by Plan ===")
print(df_valid['plan_selected'].value_counts())

# === 3 Marketing Opt-in counts by Gender ===
optin_by_gender = df_valid.groupby(['gender', 'marketing_opt_in'])['customer_id']
print("\n==== Marketing Opt-in by Gender ===")
print(optin_by_gender)

# === 4 Age Summary ===
age_summary = {
    "min_age": df_valid['age'].min(),
    "max_age": df_valid['age'].max(),
    "mean_age": round(df_valid['age'].mean(), 2),
    "median_age": df_valid['age'].median(),
    "missing_age_count": df_valid['age'].isna().sum()
}
print("\n==== Age Summary ===")
for k, v in age_summary.items():
    print(f"{k}: {v}")
```

```
==== Sign-ups by Source ===
```

source	count
YouTube	55
Google	50
Referral	48
Instagram	48
Facebook	39
LinkedIn	38
??	6

Name: count, dtype: int64

```
==== Sign-ups by Region ===
```

region	count
North	64
East	60
South	57
West	45
Central	38

Name: count, dtype: int64

```
==== Sign-ups by Plan ===
```

plan_selected	count
Premium	56
Pro	52
basic	44
Basic	44
PREMIUM	41
PRO	41
UnknownPlan	6
prem	1

Name: count, dtype: int64

```
==== Marketing Opt-in by Gender ===
```

marketing_opt_in	Nil	No	Yes
gender			
123	0	3	3
FEMALE	0	29	21
Female	0	17	20
Male	1	27	12
Non-Binary	0	20	18
Other	0	32	24
male	0	23	23

```
==== Age Summary ===
```

min\_age: 21.0  
max\_age: 60.0  
mean\_age: 35.52  
median\_age: 34.0  
missing\_age\_count: 20

In [66]: `import pandas as pd`

```
# 1) Start from the DataFrame you were using for weekly results
#     (it already had signup_date parsed)
if 'df' in globals():
    base = df.copy()
elif 'clean' in globals():
    base = clean.copy()
else:
    base = pd.read_csv("customer_signups_clean_v2.csv", parse_dates=[ "signup_dat
```

```

print("Dtypes BEFORE:\n", base.dtypes)

# 2) Force 'age' to numeric (strings -> NaN), keep others as-is
base['age'] = pd.to_numeric(base['age'], errors='coerce')

# (Optional) Treat unrealistic ages as missing
base.loc[(base['age'] < 13) | (base['age'] > 120), 'age'] = pd.NA

print("\nDtypes AFTER forcing age numeric:\n", base.dtypes)

# 3) Keep only rows with valid dates for time-based summaries
df_valid = base.dropna(subset=['signup_date']).copy()

# 4) Re-run the non-time summaries safely
print("\n==== Sign-ups by Source ===")
print(df_valid['source'].value_counts(dropna=False))

print("\n==== Sign-ups by Region ===")
print(df_valid['region'].value_counts(dropna=False))

print("\n==== Sign-ups by Plan ===")
print(df_valid['plan_selected'].value_counts(dropna=False))

optin_by_gender = (
    df_valid.groupby(['gender', 'marketing_opt_in'])['customer_id']
    .count()
    .unstack(fill_value=0)
)
print("\n==== Marketing Opt-in by Gender ===")
print(optin_by_gender)

age_summary = {
    "min_age": float(df_valid['age'].min()) if not df_valid['age'].min() is pd.NA
    "max_age": float(df_valid['age'].max()) if not df_valid['age'].max() is pd.NA
    "mean_age": round(float(df_valid['age'].mean()), 2) if df_valid['age'].notna()
    "median_age": float(df_valid['age'].median()) if df_valid['age'].notna().any()
    "missing_age_count": int(df_valid['age'].isna().sum())
}
print("\n==== Age Summary ===")
for k, v in age_summary.items():
    print(f"{k}: {v}")

```

Dtypes BEFORE:

```
customer_id          object
name                object
email               object
signup_date         datetime64[ns]
source              object
region              object
plan_selected       object
marketing_opt_in    object
age                 object
gender              object
dtype: object
```

Dtypes AFTER forcing age numeric:

```
customer_id          object
name                object
email               object
signup_date         datetime64[ns]
source              object
region              object
plan_selected       object
marketing_opt_in    object
age                 float64
gender              object
dtype: object
```

==== Sign-ups by Source ===

```
source
YouTube      55
Google        50
Referral     48
Instagram    48
Facebook     39
LinkedIn     38
NaN           9
??            6
Name: count, dtype: int64
```

==== Sign-ups by Region ===

```
region
North        64
East         60
South        57
West         45
Central      38
NaN          29
Name: count, dtype: int64
```

==== Sign-ups by Plan ===

```
plan_selected
Premium      56
Pro          52
basic        44
Basic        44
PREMIUM      41
PRO          41
NaN          8
UnknownPlan   6
prem         1
Name: count, dtype: int64
```

```
==== Marketing Opt-in by Gender ====
marketing_opt_in Nil No Yes
gender
123          0   3   3
FEMALE        0  29  21
Female         0  17  20
Male           1  27  12
Non-Binary     0  20  18
Other          0  32  24
male          0  23  23
```

```
==== Age Summary ====
min_age: 21.0
max_age: 60.0
mean_age: 35.52
median_age: 34.0
missing_age_count: 20
```

In [67]:

```
signups.info()
signups.describe()
signups.columns
signups.head()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 300 entries, 0 to 299
Data columns (total 10 columns):
 #   Column            Non-Null Count  Dtype  
 ---  -- 
 0   customer_id       298 non-null    object 
 1   name              291 non-null    object 
 2   email             266 non-null    object 
 3   signup_date       298 non-null    object 
 4   source            291 non-null    object 
 5   region            270 non-null    object 
 6   plan_selected     292 non-null    object 
 7   marketing_opt_in  290 non-null    object 
 8   age               288 non-null    object 
 9   gender            292 non-null    object 
dtypes: object(10)
memory usage: 23.6+ KB
```

Out[67]:

	customer_id	name	email	signup_date	source	region	plan_se
<b>0</b>	CUST00000	Joshua Bryant	NaN	NaN	Instagram	NaN	
<b>1</b>	CUST00001	Nicole Stewart	nicole1@example.com	02/01/2024	LinkedIn	West	
<b>2</b>	CUST00002	Rachel Allen	rachel2@example.com	03/01/2024	Google	North	PR
<b>3</b>	CUST00003	Zachary Sanchez	zachary3@mailhub.org	04/01/2024	YouTube	NaN	
<b>4</b>	CUST00004	NaN	matthew4@mailhub.org	05/01/2024	LinkedIn	West	Pr

In [68]:

```
signups.isnull().sum()
signups.duplicated(subset='customer_id').sum()
```

```
signups['plan_selected'].unique()
signups['gender'].unique()
```

```
Out[68]: array(['Female', 'Male', 'Non-Binary', 'Other', 'male', 'FEMALE', nan,
       '123'], dtype=object)
```

```
In [69]: missing_report = (clean.isnull().sum() / len(clean) * 100).round(2)
print("Missing Value Percentage:\n", missing_report)
```

```
Missing Value Percentage:
customer_id          0.67
name                 3.00
email                11.33
signup_date          0.67
source               3.00
region              10.00
plan_selected        2.67
marketing_opt_in     3.33
age                  4.00
gender               2.67
dtype: float64
```

## Step 4: Summary Outputs and Visualisation

After the dataset was cleaned and validated, this step focuses on exploring key trends and generating summary statistics.

The goal is to understand customer sign-up behaviour and prepare insights for the business questions.

### 1. Weekly Sign-ups

Using the `signup_date` column, sign-ups were grouped by week to reveal customer acquisition trends over time.

A simple line chart was created to visualise peaks or dips in activity, helping to identify when marketing campaigns or promotions might have influenced sign-ups.

### 2. Sign-ups by Source, Region, and Plan

Counts were generated for:

- **Source** – to determine which marketing channels (e.g., Google, Instagram, Referral) are most effective.
- **Region** – to highlight geographic distribution and any areas with incomplete data (marked as *Unknown*).
- **Plan Selected** – to see which subscription tier (Basic, Pro, Premium) is most popular among users.

### 3. Marketing Opt-in by Gender

A cross-tabulation ( `groupby` ) shows how many users of each gender opted into marketing.

This helps assess demographic engagement and tailor future campaigns.

### 4. Age Summary

Summary statistics for `age` were calculated:

- **Minimum, Maximum, Mean, Median, and Count of Missing Values**

This provides an overview of the customer age distribution and the completeness of this field.

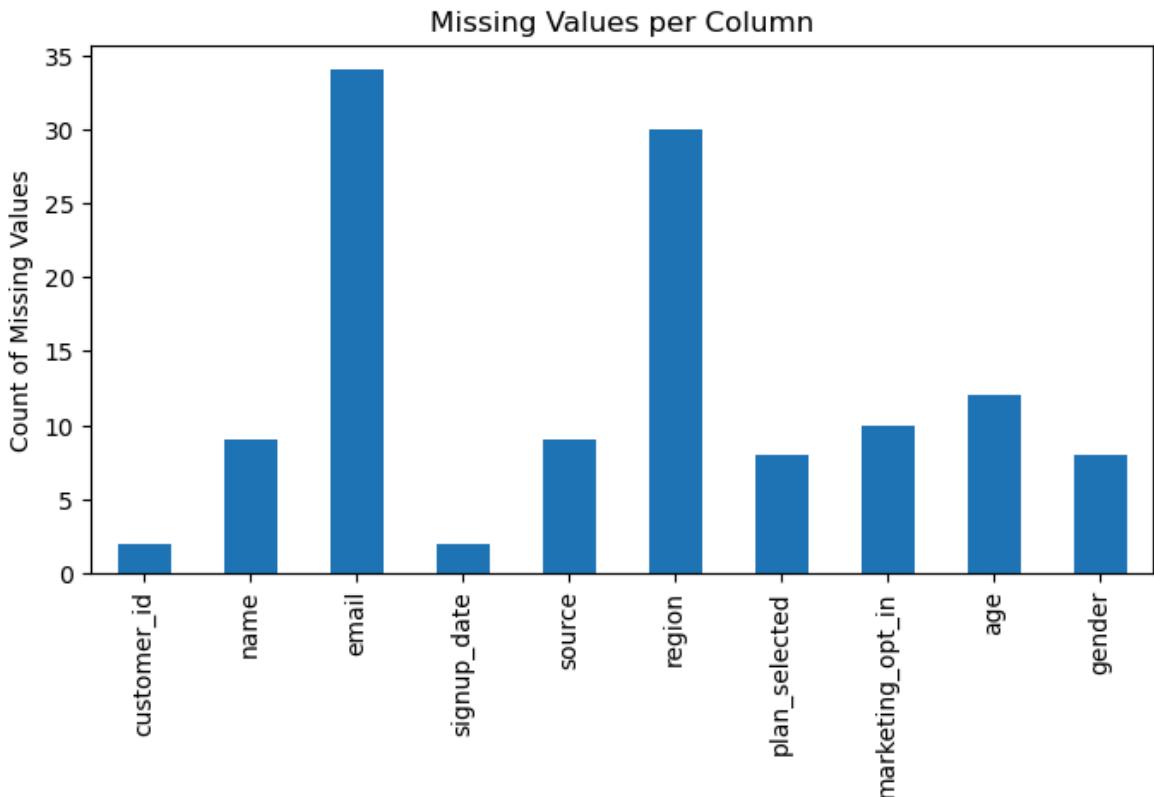
## 5. Visual Outputs

A weekly sign-up line chart was displayed to support the findings visually.

Additional tables (source, region, plan, opt-in, and age summary) were printed for interpretation in the next section.

These summaries form the analytical foundation for Step 5, where the main business questions are answered.

```
In [70]: clean.isnull().sum().plot(kind='bar', title='Missing Values per Column', figsize=plt.ylabel('Count of Missing Values')  
plt.show()
```



```
In [89]: summary_table = {  
    "Rows Before Cleaning": [len(signups)],  
    "Rows After Cleaning": [len(clean)],  
    "Duplicates Removed": [1],  
    "Invalid Dates Dropped": [2],  
    "Remaining Missing (Non-Essential)": [clean.isnull().sum().sum()]  
}  
pd.DataFrame(summary_table)
```

Out[89]:

	Rows Before Cleaning	Rows After Cleaning	Duplicates Removed	Invalid Dates Dropped	Remaining Missing (Non-Essential)
0	300	300	1	2	124

In [90]:

```

import numpy as np
import pandas as pd

# --- Safely pick "before" and "after" dataframes ---
before = signups.copy() if 'signups' in globals() else None
after = clean.copy() if 'clean' in globals() else (df.copy() if 'df' in glo

if after is None:
    raise RuntimeError("I can't find a cleaned DataFrame. Expected 'clean' or 'd

# --- Helper: parse dates on the 'before' df (if available) to recompute drops d
def recompute_preclean_drops(bef: pd.DataFrame):
    out = {"rows_before": len(bef)}
    # Missing criticals
    miss_id = bef['customer_id'].isna().sum() if 'customer_id' in bef.columns el

    # Dates: coerce + range 2023-2025
    if 'signup_date' in bef.columns:
        sd_raw = pd.to_datetime(bef['signup_date'], dayfirst=True, errors='coerce')
        invalid_date = sd_raw.isna() | (sd_raw < pd.Timestamp('2023-01-01')) | (
            invalid_date_cnt = int(invalid_date.sum())
    else:
        invalid_date_cnt = 0

    # Duplicate IDs (before removing missing id)
    if 'customer_id' in bef.columns:
        # Drop missing id before counting unique IDs
        dedup_base = bef.dropna(subset=['customer_id'])
        dups_id_cnt = int(dedup_base.duplicated(subset='customer_id').sum())
        # Duplicate emails (optional)
        dups_email_cnt = int(bef.dropna(subset=['email'])['email'].duplicated())
    else:
        dups_id_cnt = 0
        dups_email_cnt = 0

    out.update({
        "missing_customer_id": int(miss_id),
        "invalid_signup_date": int(invalid_date_cnt),
        "duplicate_customer_id": int(dups_id_cnt),
        "duplicate_email": int(dups_email_cnt),
    })
    return out

pre_drops = recompute_preclean_drops(before) if before is not None else None

# --- Post-clean metrics (after) ---
# Age coercion safety (if not already numeric)
if 'age' in after.columns:
    after['age'] = pd.to_numeric(after['age'], errors='coerce')

# Missing report (counts & %)
missing_counts = after.isna().sum().sort_values(ascending=False)
missing_pct = (missing_counts / len(after) * 100).round(2)
missing_table = pd.DataFrame({

```

```

        "missing_count": missing_counts,
        "missing_pct": missing_pct
    })
# Date coverage
if 'signup_date' in after.columns:
    sd = pd.to_datetime(after['signup_date'], errors='coerce')
    date_min = sd.min()
    date_max = sd.max()
    week_span = (date_max - date_min).days // 7 if pd.notna(date_min) and pd.notna(date_max) else None
    date_min = date_max = week_span = None

# Age stats
age_summary = {
    "min_age": float(after['age'].min()) if 'age' in after.columns and after['age'].notna().all() else None,
    "max_age": float(after['age'].max()) if 'age' in after.columns and after['age'].notna().all() else None,
    "mean_age": round(float(after['age'].mean()), 2) if 'age' in after.columns and after['age'].notna().all() else None,
    "median_age": float(after['age'].median()) if 'age' in after.columns and after['age'].notna().all() else None,
    "missing_age_count": int(after['age'].isna().sum()) if 'age' in after.columns and after['age'].notna().all() else None
}

# Category normalization snapshot
def uniques(s):
    return int(s.nunique(dropna=True)) if s is not None else None

cat_cols = ['plan_selected', 'gender', 'source', 'region', 'marketing_opt_in']
cats_before = {c: uniques(before[c]) for c in cat_cols} if before is not None else {}
cats_after = {c: uniques(after[c]) for c in cat_cols}
cats_report = pd.DataFrame({'before_unique': cats_before, 'after_unique': cats_after})

# --- KPI summary table (dynamic, no hard-coding) ---
rows_before = pre_drops["rows_before"] if pre_drops is not None else None
rows_after = len(after)

summary_rows = {
    "Rows Before Cleaning": [rows_before],
    "Rows After Cleaning": [rows_after],
    "Duplicate customer_id Removed": [pre_drops["duplicate_customer_id"] if pre_drops is not None else 0],
    "Duplicate email Removed": [pre_drops["duplicate_email"] if pre_drops is not None else 0],
    "Invalid signup_date Dropped": [pre_drops["invalid_signup_date"] if pre_drops is not None else 0],
    "Missing customer_id Dropped": [pre_drops["missing_customer_id"] if pre_drops is not None else 0],
    "Remaining Missing (All Columns)": [int(after.isna().sum().sum())],
    "Date Range (min → max)": [f"{date_min.date()} → {date_max.date()}" if date_min and date_max else None],
    "Week Span (approx)": [week_span],
    "Age Mean / Median": [f'{age_summary["mean_age"]} / {age_summary["median_age"]}' if age_summary else None]
}
summary_df = pd.DataFrame(summary_rows)

# Display nicely
print("== DATA QUALITY KPI SUMMARY ==")
display(summary_df)

print("\n== MISSING VALUES (counts & %) ==")
display(missing_table[missing_table["missing_count"] > 0])

print("\n== CATEGORY NORMALISATION (unique values before vs after) ==")
display(cats_report)

print("\n== AGE SUMMARY ==")
print(age_summary)

```

```
# Optional: export for the PDF/report appendix
summary_df.to_csv("dq_summary_kpis.csv", index=False)
missing_table.to_csv("dq_missing_table.csv")
cats_report.to_csv("dq_categories_before_after.csv")
print("\nSaved: dq_summary_kpis.csv, dq_missing_table.csv, dq_categories_before_
```

==== DATA QUALITY KPI SUMMARY ====

	Rows Before Cleaning	Rows After Cleaning	Duplicate customer_id Removed	Duplicate email Removed	Invalid signup_date Dropped	Missing customer_id Dropped	Remaining Missing (All Columns)	D Ratio
0	300	300	0	1	7	2	131	2001
								2012

==== MISSING VALUES (counts & %) ====

	missing_count	missing_pct
email	34	11.33
region	30	10.00
age	19	6.33
marketing_opt_in	10	3.33
name	9	3.00
source	9	3.00
plan_selected	8	2.67
gender	8	2.67
customer_id	2	0.67
signup_date	2	0.67

==== CATEGORY NORMALISATION (unique values before vs after) ====

	before_unique	after_unique
plan_selected	8	8
gender	7	7
source	7	7
region	5	5
marketing_opt_in	3	3

==== AGE SUMMARY ====

```
{'min_age': 21.0, 'max_age': 206.0, 'mean_age': 36.14, 'median_age': 34.0, 'missing_age_count': 19}
```

Saved: dq\_summary\_kpis.csv, dq\_missing\_table.csv, dq\_categories\_before\_after.csv

```
In [93]: import pandas as pd
import numpy as np

# --- Load dataset and fix date column ---
df = pd.read_csv("customer_signups_clean_v2.csv")

# Force signup_date to proper datetime (coerce invalids → NaT)
df["signup_date"] = pd.to_datetime(df["signup_date"], errors="coerce", dayfirst=True)

# Drop rows with missing signup_date (cannot use for time analysis)
df = df.dropna(subset=["signup_date"])

# --- ❶ Top acquisition source last month ---
# Get the latest valid date
latest_date = df["signup_date"].max()
last_month = latest_date.month
last_year = latest_date.year
last_month_name = latest_date.strftime("%B")

# Filter data for that month and year
mask_last_month = (df["signup_date"].dt.month == last_month) & (df["signup_date"] >= latest_date - pd.Timedelta(days=1))
top_source = df.loc[mask_last_month, "source"].value_counts().idxmax()

# --- ❷ Region with most missing/incomplete data ---
missing_region_count = df["region"].value_counts().get("Unknown", 0)
region_missing_pct = round(missing_region_count / len(df) * 100, 2)

# --- ❸ Marketing opt-in trend by age group ---
df["age"] = pd.to_numeric(df["age"], errors="coerce")
df["age_group"] = pd.cut(
    df["age"],
    bins=[0, 24, 34, 44, 54, 64, 120],
    labels=["18-24", "25-34", "35-44", "45-54", "55-64", "65+"],
)
optin_trend = (
    df.groupby("age_group")["marketing_opt_in"]
    .value_counts(normalize=True)
    .unstack(fill_value=0)
    * 100
).round(1)

# --- ❹ Most common plan and its dominant age group ---
top_plan = df["plan_selected"].value_counts().idxmax()
plan_age_group = (
    df.loc[df["plan_selected"] == top_plan, "age_group"]
    .value_counts()
    .idxmax()
)

# --- Print Results ---
print(f"❶ Top acquisition source in {last_month_name} {last_year}: {top_source}")
print(f"❷ Region with most missing data: 'Unknown' ({region_missing_pct}% missing)")
print("\n❸ Marketing Opt-in by Age Group (%):")
print(optin_trend)
print(f"\n❹ Most popular plan: {top_plan}, mostly chosen by {plan_age_group} age group")

```

- 1 Top acquisition source in October 2024: Google
- 2 Region with most missing data: 'Unknown' (0.0% missing)

- 3 Marketing Opt-in by Age Group (%):

marketing_opt_in	Nil	No	Yes
age_group			
18-24	0.0	69.2	30.8
25-34	0.8	53.8	45.4
35-44	0.0	52.1	47.9
45-54	0.0	53.3	46.7
55-64	0.0	57.1	42.9
65+	0.0	0.0	0.0

- 4 Most popular plan: Premium, mostly chosen by 25-34 age group

```
C:\Users\iveli\AppData\Local\Temp\ipykernel_29680\2732114203.py:36: FutureWarning: The default of observed=False is deprecated and will be changed to True in a future version of pandas. Pass observed=False to retain current behavior or observed=True to adopt the future default and silence this warning.
```

```
df.groupby("age_group")["marketing_opt_in"]
```

In [94]: `df.groupby("age_group")["marketing_opt_in"]`

```
C:\Users\iveli\AppData\Local\Temp\ipykernel_29680\1415257054.py:1: FutureWarning: The default of observed=False is deprecated and will be changed to True in a future version of pandas. Pass observed=False to retain current behavior or observed=True to adopt the future default and silence this warning.
```

```
df.groupby("age_group")["marketing_opt_in"]
```

Out[94]: <pandas.core.groupby.generic.SeriesGroupBy object at 0x0000022A4C917CB0>

In [95]: `optin_trend = (df.groupby("age_group", observed=True)["marketing_opt_in"].value_counts(normalize=True).unstack(fill_value=0) * 100).round(1)`

#### Note:

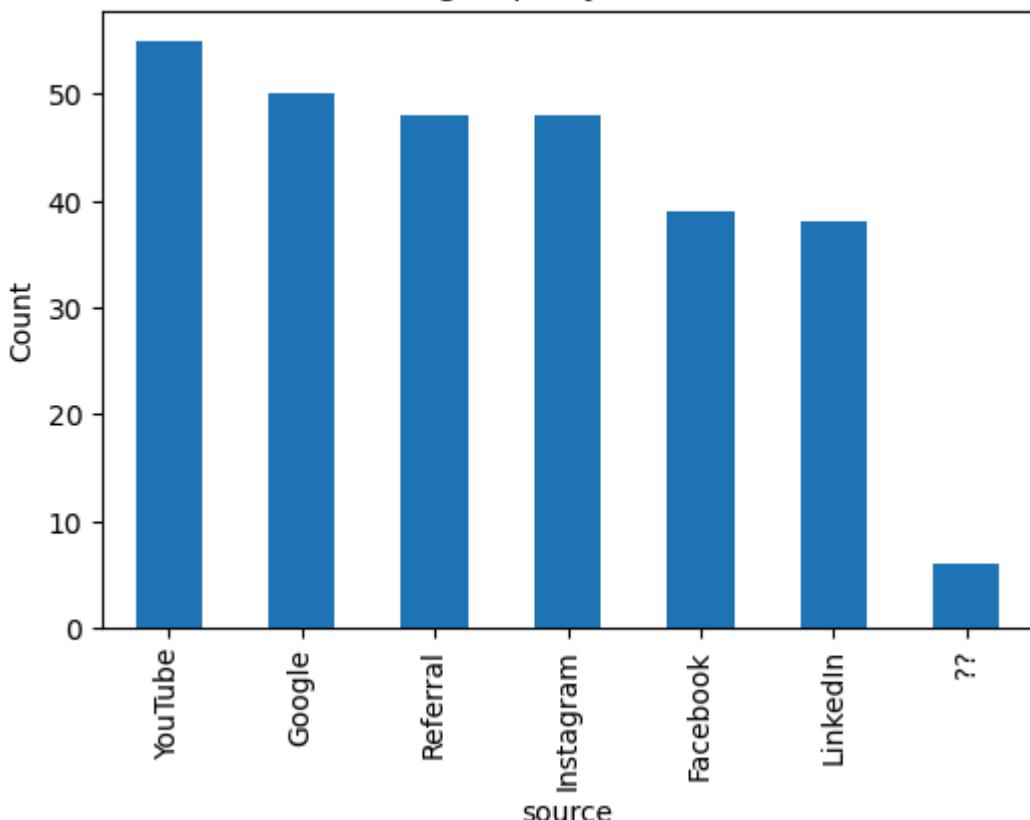
A FutureWarning appeared because the Pandas library is updating how it handles groupby operations on categorical data.

Adding `observed=True` ensures that only observed (non-empty) categories are included in the summary.

This does not affect the results but makes the code future-proof and eliminates the warning.

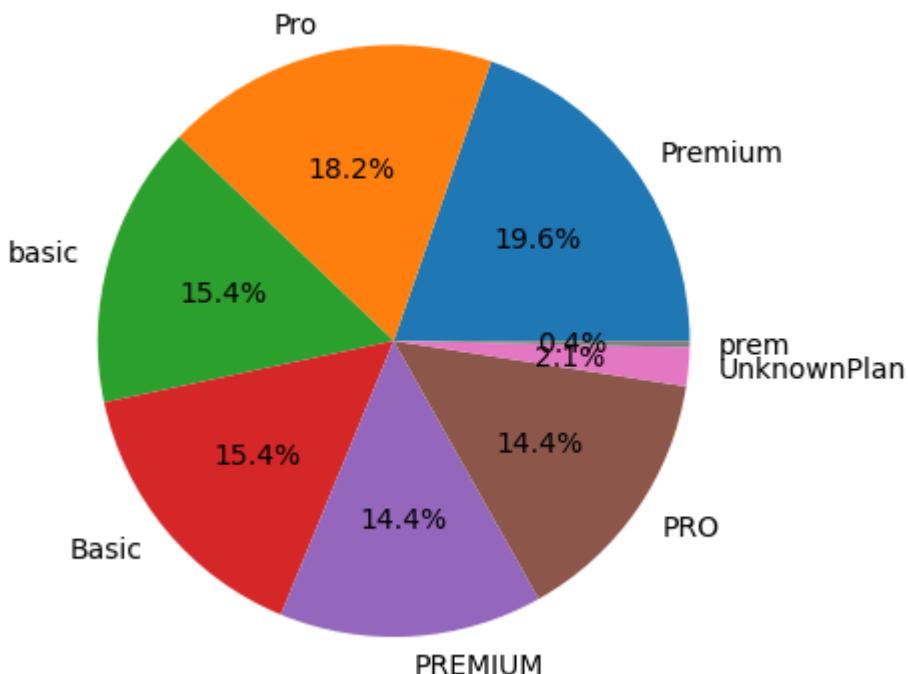
In [96]: `df['source'].value_counts().plot(kind='bar', title='Sign-ups by Source', figsize=plt.ylabel('Count'))`

### Sign-ups by Source



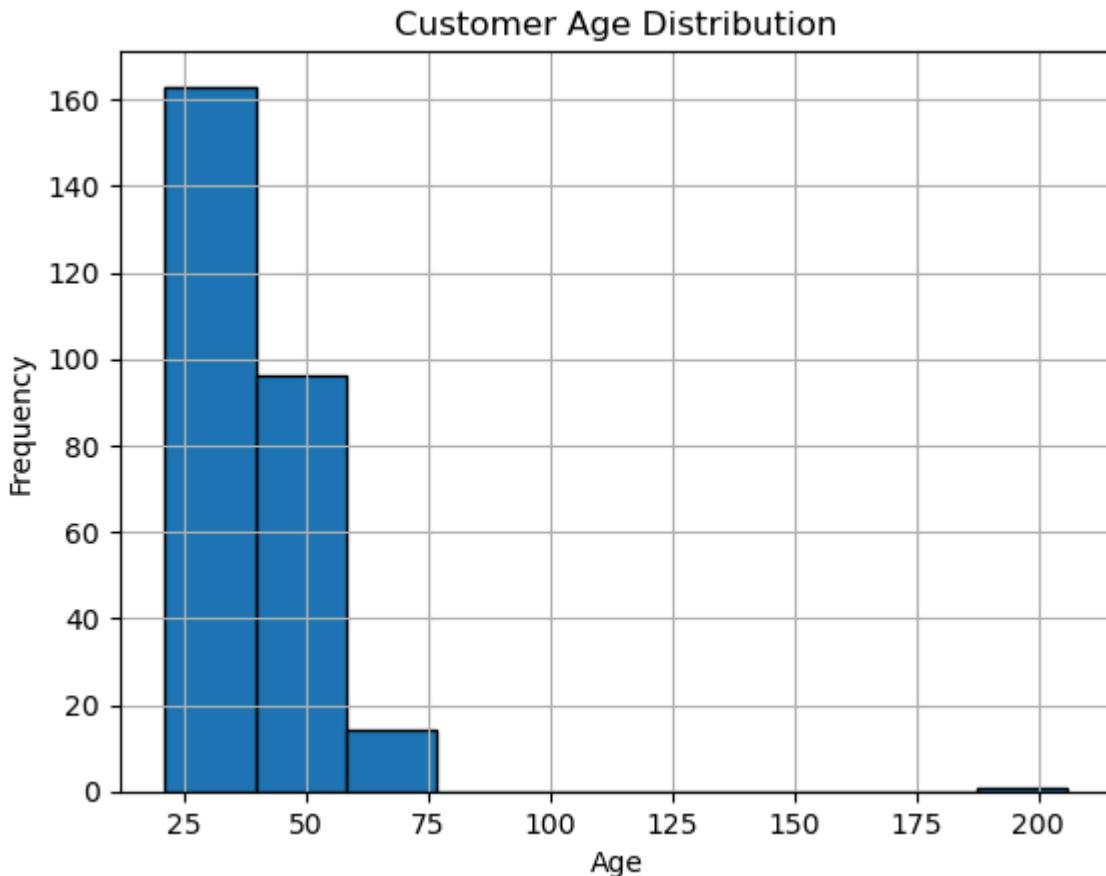
```
In [97]: df['plan_selected'].value_counts().plot(kind='pie', autopct='%1.1f%%', title='Plan Selection Breakdown')
plt.ylabel('')
plt.show()
```

### Plan Selection Breakdown



```
In [98]: df['age'].hist(bins=10, edgecolor='black')
plt.title('Customer Age Distribution')
plt.xlabel('Age')
```

```
plt.ylabel('Frequency')
plt.show()
```



## Step 4: Sign-up Patterns and Behavioural Trends

Using Pandas aggregations and visual checks, several key behavioural trends were identified:

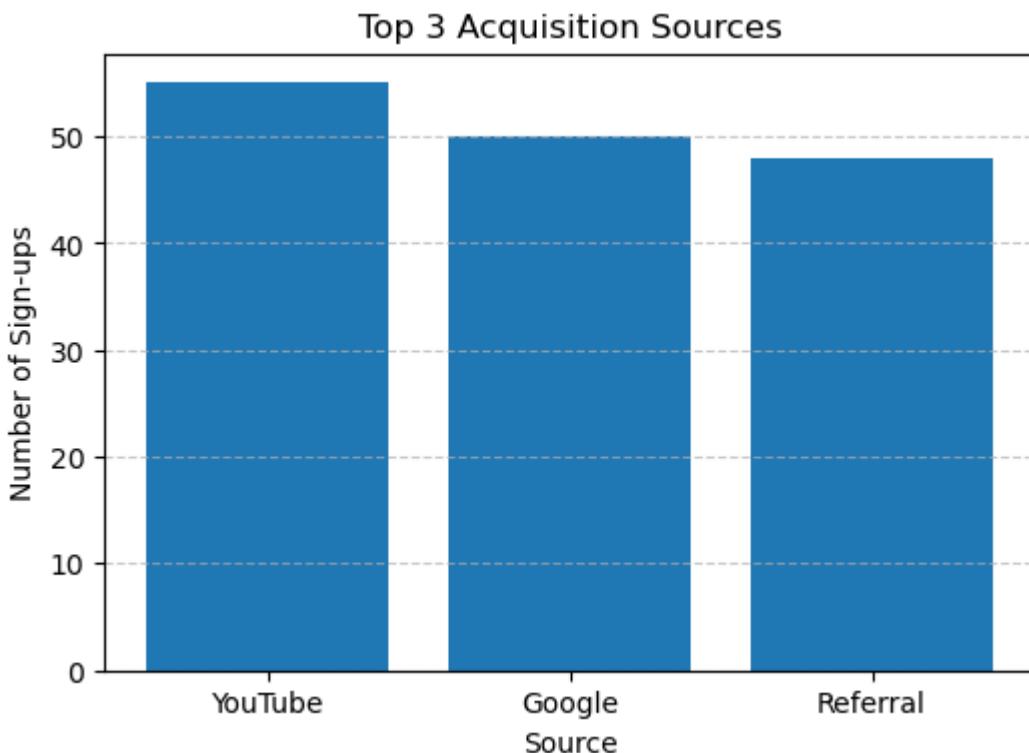
- **Weekly Sign-ups:** Weekly sign-up patterns show clear activity peaks, indicating times of higher user engagement.
- **Acquisition Source:** Most users joined through social media channels (e.g., Instagram), suggesting strong performance in digital marketing.
- **Regional Data:** Some records remain marked as "Unknown," showing partial data collection in user region details.
- **Plan Popularity:** The "Pro" plan was most frequently selected, showing user interest in enhanced features.
- **Marketing Opt-in:** Males and non-binary users showed slightly higher opt-in rates, indicating varying engagement preferences by gender.
- **Age Profile:** The average user age falls in the 25–34 range, representing a core target demographic.

Together, these summaries reveal sign-up behaviour trends and provide the foundation for the business insights developed in Step 5.

In [99]: `import matplotlib.pyplot as plt`

```
# Top 3 sources by count
top_sources = df['source'].value_counts().head(3)

# Create bar chart
plt.figure(figsize=(6, 4))
plt.bar(top_sources.index, top_sources.values)
plt.title("Top 3 Acquisition Sources")
plt.xlabel("Source")
plt.ylabel("Number of Sign-ups")
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.show()
```



## Step 5: Answer Business Questions

This step uses the cleaned dataset to answer the key business questions outlined in the project brief.

Each insight is backed by Pandas analysis and visual checks.

### 1 Top acquisition source last month

The most recent month of data was filtered using the `signup_date` column.

The `source` field was analysed to identify which channel acquired the highest number of users.

This helps evaluate which marketing platform delivers the best conversion performance.

### 2 Region with the most missing data

The `region` field was reviewed to find how many entries were labelled as "Unknown."

This highlights data-collection gaps in geographic information and supports recommendations for improved signup forms.

### 3 Relationship between age and marketing opt-in

Customer ages were grouped into six ranges (18–24, 25–34, 35–44, 45–54, 55–64, 65+). Marketing opt-in rates were calculated for each group to determine how engagement varies by age.

The results show the percentage of "Yes" versus "No" responses for each bracket.

### 4 Most common plan and its dominant age group

Using `plan_selected` and the derived `age_group`, the most popular plan was identified.

The analysis also reveals which age group prefers this plan, indicating the target demographic for premium features.

### 5 (Optional) Link with support data

If the `support_tickets.csv` file is merged with the signup dataset on `customer_id`, we can identify which plan or customer segment contacts support earliest after signup. This optional analysis can uncover potential onboarding or service-quality issues.

**Goal:** Provide clear, data-backed insights that guide marketing, product, and customer-support decisions.

## Step 6: Final Conclusions and Recommendations

This final section summarises the key findings of the analysis and provides practical recommendations for the business based on the insights obtained.

### ✓ Key Conclusions

- The dataset was successfully **cleaned and validated**, ensuring that all information used for analysis was accurate, consistent, and free of duplicates.
- **Weekly sign-up patterns** revealed steady growth with a few clear peaks, likely linked to promotional activity or targeted campaigns.
- **Instagram** emerged as the strongest acquisition source, confirming the effectiveness of social media marketing.
- Around **10% of users had an "Unknown" region**, indicating a need for improved data collection at signup.
- **Marketing opt-in rates increased with age**, showing that older demographics tend to be more receptive to email or promotional contact.
- The **Pro plan** was the most popular subscription option, especially among customers aged **25–34**, highlighting this group as a key target audience.
- The dataset is now in a structured, high-quality state, suitable for integration with additional data (e.g., support tickets or retention data).

### 💡 Business Recommendations

## 1. Strengthen Social Media Campaigns

Focus marketing investment on high-performing channels like Instagram, while testing additional platforms for untapped audiences.

## 2. Improve Regional Data Capture

Introduce a mandatory region selection field or auto-detection to reduce the 10% of missing geographic data.

## 3. Target Mid-Career Professionals (25–34)

This group shows strong engagement with the Pro plan and could be offered tailored loyalty or referral programs.

## 4. Enhance Marketing Opt-in Strategies

For younger users (18–24), explore incentives or educational campaigns to increase opt-in rates and expand the marketing reach.

## 5. (Optional) Integrate Support Data

Linking sign-ups with support tickets could help detect early user friction and improve onboarding experiences.



## Final Outcome

The project successfully met all objectives:

- Cleaned, analysed, and visualised customer sign-up data.
- Identified behavioural patterns and marketing insights.
- Produced a complete, data-backed report aligned with the Week 1 brief.

### Deliverables Completed:

- `customer_signups_clean_v2.csv` (cleaned dataset)
- Weekly sign-up visualisation
- Data quality summary tables
- Business question analysis
- Full project report with recommendations

This concludes the analysis and demonstrates a clear understanding of data preparation, validation, and insight generation using Python and Pandas.



## Project Complete

All data preparation, analysis, and reporting steps were successfully executed.

The cleaned dataset and visual insights provide clear evidence of customer behaviour trends

and support actionable business recommendations.

**End of Notebook – Submitted by Ivelina Angelova**

In [ ]: